

Anonymous Project Spec for Peer Pricing Review

Executive summary and purpose of this document

I'm trying to triangulate a fixed fee for a phased B2B consulting engagement, and I'd like an unbiased number from someone whose work I respect.

I've stripped client identity, my own firm's identity, and any signal of existing pricing. Please answer based on scope, role mix, and effort, not on what you think I want to hear or what you think I currently charge.

Alongside any feedback, there are three structured questions for you at the bottom.

Thank you. I am grateful for your help, and I'll share back what I learn.

Executive summary and purpose of this document	1
Hospitality glossary	2
Engagement at a glance	2
Client profile	2
Discipline	3
Phased scope	3
Effort and timing	4
Role mix and seniority	4
What's NOT in scope	5
Other context relevant to pricing	5
Differentiating factors that may matter for the price	5
What I'd like from you	6

Hospitality glossary

A small glossary of hospitality terms used below:

- **OTA** is Online Travel Agency, meaning Booking.com, Expedia, and similar third-party booking platforms that charge hotels 15-25% commission on bookings.
- **ADR** is Average Daily Rate, the average price per occupied room per night.
- **Upper-upscale** is the second-highest tier in the standard hotel-industry classification (luxury, upper-upscale, upscale, upper-midscale, midscale, economy), roughly equivalent to four-and-a-half-star positioning.

Engagement at a glance

- A focused, four-phase consulting project.
- Discovery, specification, coordination of implementation, and validation.
- Technical implementation is handled by the client's existing technical partner, not by the consultancy.
- The consultancy designs, specs, supports, and validates.

Client profile

- European boutique hospitality property in the upper-upscale tier (industry tier classification, see glossary above), family-owned, owner-operator decision-maker.
- About 50 rooms, niche wellness positioning in a regional specialty.
- Member of a major international hotel collection (premium brand affiliation that signals positioning but leaves marketing and website ownership to the property).
- Recently completed major renovation. Live website, recent translation work just shipped.
- Three personas in the buying conversation:
 - the owner-operator (cares about bookings, ADR, OTA-independence, ROI)
 - the in-house brand marketer (not technical, brand-and-feeling led)
 - and the property's external technical implementation partner (web-aware, can write code, will execute the spec).
- Existing relationship with the consultancy from a prior, successful, smaller engagement.

Discipline

A discrete piece of B2B strategy and technical content work at the intersection of:

- Information architecture and structured-data design (schema.org spec)
- AI-search visibility and generative-engine-optimization (GEO/AEO) advisory
- Content design, specifically writing schema-block content values in **two languages** (DE + EN)
- Coordination with the client's technical implementation partner on architectural choices, async QA, and deploy-time verification
- Validation against an agreed pre/post AI-engine query baseline

The deliverable is

- a complete, page-by-page schema specification
- an AI-engine optimization layer (including llms.txt and llms-full.txt recommendations)
- and a comparison report at the end

Phased scope

Phase 0: Baseline capture (~3 senior consultant days)

- Inventory and validation of existing structured data (what's there, what's valid, what's missing)
- Google Search Console snapshot for an agreed query set
- Organic ranking baseline for ~10-15 segment-relevant queries
- Documented behavior of 5 AI engines (ChatGPT, Perplexity, Claude, Gemini, Google AI Overviews) on 10-15 segment queries, with screenshots and transcripts
- Short baseline write-up for comparison in Phase 3

Phase 1: Schema design and specification (~8 senior consultant days)

- Schema architecture across ~6-7 page-type variants (Hotel root, room/lodging-unit, restaurant, spa/wellness, FAQ, event, breadcrumb enrichment)
- Bilingual content values (DE + EN) for every schema block; every block has text copy that needs drafting in two languages
- AI-engine optimization layer: research-backed best practice translation, **llms.txt** and **llms-full.txt** drafting, citation-friendly content recommendations
- Page-by-page specification document, JSON-LD blocks ready to implement, or extension configuration spec (architectural choice deferred to the client's technical partner during this phase)
- Coordination with the technical partner on architectural route selection

Phase 2: Implementation coordination (~2 senior consultant days)

- Architectural-route discussion and decisions with the client's technical partner
- Async QA support during implementation, including clarifications and partial reviews
- Final QA of deployed markup against the spec; validation pass DE + EN, sample checks across page types

Phase 3: Validation and handover (~2 senior consultant days)

- Re-run the same AI-engine query set from Phase 0 against the live, marked-up site, using the same engines and the same queries, with screenshots and transcripts (3-4 months after Phase 0)
- Rich Results Test and Schema.org Validator full-site pass (DE + EN)
- Comparison report against Phase 0 baseline
- Living-spec handover documentation that the client can extend to future pages

Effort and timing

- **Total senior consultant effort:** ~15 days across the four phases (Phase 0: 3, Phase 1: 8, Phase 2: 2, Phase 3: 2).
- **Calendar duration:** ~3-4 months elapsed end to end, with most consultancy effort concentrated in Phase 1 and the rest of the calendar driven by client decisions, the technical partner's implementation timeline, and post-deploy stabilization between Phase 2 and Phase 3.
- **Distribution of effort across the calendar:** front-loaded around weeks 1-6, then thin advisory and QA touchpoints across weeks 7-12, then a concentrated validation push at weeks 13-16.

Role mix and seniority

- **Single senior consultant** delivering across all phases, including strategy, persona work, schema architecture, content drafting, technical content, and validation. No junior pyramid. No subcontracted research.
- The consultancy is owner-operator scale, boutique, narrow specialization. Not enterprise strategy, not generalist content marketing.
- Bilingual delivery (DE + EN) is performed by the same senior person, not outsourced.

What's NOT in scope

Handled by the client's technical partner, separately commissioned by them:

- Actual implementation of the spec in the client's CMS (Symfony / Contao stack)
- Hosting, deployment, server-side performance work
- Front-end development
- Ongoing site maintenance after handover

The consultancy's role is design, specification, advisory, and validation, not execution. The technical partner's fee is between them and the client.

Other context relevant to pricing

- Client is in EU; consultancy bills net of VAT under intra-EU B2B reverse-charge.
- Currency: EUR.
- Payment structure: 25% deposit on Phase 0 kickoff, then 25% on each phase completion (even split across the four phases). Net 14.
- Engagement model: fixed fee per phase, sign-off gate at each phase boundary, scope changes between phases handled as separate addenda.
- The client is mid-conversation, not a cold pitch; the prior engagement is already a known-good signal of the consultancy's quality.
- The consultancy has been undercharging on prior work with this same client (a deliberately under-priced first engagement) and is recalibrating going forward.

Differentiating factors that may matter for the price

- **The schema/GEO and AI-search-readiness combination, executed bilingually, is a niche specialization.**
- An audit of 7 peer properties in the same hotel collection (and the collection's own brand site) found that none currently emit the entity-graph schema relationship the spec will create.
 - →→→ **So the deliverable is a measurable competitive differentiator for the client, not a parity move.**
- The validation phase produces evidence (specifically, a pre/post AI-engine query comparison segmented by language) that the client can use internally and with the parent brand collection.

What I'd like from you

Three short questions. Please commit to a number on each. The most useful thing for me is your gut estimate, not a hedge.

1. What fixed-fee total would you propose for this scope? Single number or a tight range. If you'd structure it differently (for example, retainer, day-rate-against-cap, value-based), say so and give the equivalent.

2. What day rate would you be applying for senior boutique consulting work like this, in DACH and/or US? The number you'd actually charge, not a market average. I'm trying to calibrate against what serious boutique operators are billing right now.

3. If you wouldn't take this engagement at all, what would change your mind? Scope cut? Different client tier? Different deliverable shape? I'm trying to understand whether the "no" signal would be price, scope, fit, or something else.

Optional bonus: any other things about the spec that would change your number, up or down, that I haven't surfaced. Push back on anything that smells off.